

PHYSICS-INFORMED NEURAL NETWORKS FOR REMAINING-USEFUL-LIFE PREDICTION OF ROLLING-ELEMENT BEARINGS



A thesis submitted in partial fulfillment of
requirements for the degree of

Bachelor of Science

in

Mechanical Engineering

Submitted By:

Mahdee Hassan Sami

Student ID: 2003101

Supervised By:

Dr. Md. Abu Mowazzem Hossain

Professor

Department of Mechanical Engineering

**CHITTAGONG UNIVERSITY OF ENGINEERING
AND TECHNOLOGY, CHATTOGRAM**

(PROJECT AND THESIS)

COURSE NO: ME-498

Title of Project: Physics-Informed Neural Networks for Remaining-Useful-Life Prediction of Rolling-Element Bearings

Name of the Student: Mahdee Hassan Sami

Student ID: 2003101

Session: 2024-2025

Name of the Department: Mechanical Engineering

Program: B.Sc. Engg.

Name of the Supervisor: Dr. Md. Abu Mowazzem Hossain
Professor
Department of Mechanical Engineering, CUET

Signature of the Student:

Signature of the Supervisor:

Abstract

Rolling-element bearings are small components, but their condition often decides whether a rotating machine can keep running safely. A bearing fault can begin as a local surface defect and then grow into severe vibration, heat, loss of alignment, secondary damage, and unplanned shutdown. Remaining useful life prediction is therefore a central problem in prognostics and health management. This thesis studies bearing remaining useful life prediction using the NASA IMS run-to-failure bearing data and compares a physics-informed neural network against three data-driven baselines. The original proposal for this work expected a physics-informed model to outperform a purely data-driven model. The completed local experiments gave a more measured result. The proposed DeepXDE physics-informed model performed best in one of the three cross-bearing experiments, but the LSTM baseline gave the lowest average RMSE and highest average R^2 across the full experiment set. This does not make the physics-informed approach unimportant. It shows that weak physical regularization, when applied to noisy run-to-failure vibration features, must be designed and weighted carefully. A simple monotonicity and fault-frequency prior can help in some transfer settings, but it may also restrict the network when the degradation path of the test bearing differs from the training bearings. The final pipeline reads the IMS data locally from extracted folders, removes the dependency on Google Drive, extracts time-domain and envelope spectral features, constructs normalized RUL targets, builds a PCA-based health indicator, and evaluates four models: a feed-forward neural baseline, the proposed DeepXDE physics-informed model, an LSTM sequence model, and a 1D CNN sequence model. The experiments use three train-validation-test splits designed to test cross-bearing and cross-dataset transfer. The best average RMSE was achieved by the LSTM baseline at 0.162, followed by the data-only feed-forward baseline at 0.185, the proposed physics-informed model at 0.223, and the CNN baseline at 0.235. The work contributes a reproducible local project structure, a documented feature and evaluation pipeline, and an honest analysis of where the investigated physics-informed formulation helped and where it did not.

Keywords: remaining useful life, rolling-element bearing, prognostics and health management, physics-informed neural network, IMS bearing dataset, vibration analysis, LSTM, CNN.

Acknowledgement

First, I offer my gratitude to Allah for giving me the strength, patience, and opportunity to complete this thesis work. I am deeply grateful to my undergraduate thesis supervisor, Dr. Md. Abu Mowazzem Hossain, Professor, Department of Mechanical Engineering, Chittagong University of Engineering & Technology, for his supervision, advice, and encouragement throughout this project.

I also thank my family for their constant support during the long process of learning the theory, writing code, running experiments, correcting mistakes, and preparing this thesis. Their patience made the work easier to continue when the results did not come out as simply as expected.

Finally, I acknowledge the researchers who made the IMS bearing dataset publicly available. Open run-to-failure datasets make it possible for students and researchers to test prognostic methods in a reproducible way.

Contents

Abstract	i
Acknowledgement	ii
Contents	vi
List of Figures	vii
List of Tables	ix
Nomenclature	x
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Objectives of the Study	2
1.4 Scope	3
1.5 Importance of the Study	3
1.6 Limitations	3
1.7 Thesis Outline	4
2 Literature Review	5
2.1 Introduction	5
2.2 Literature Survey	5

2.2.1	Rolling-Element Bearings: Fundamentals and Failure Modes . . .	5
2.2.2	Traditional Predictive Maintenance Techniques for Bearings . . .	6
2.2.3	Deep Learning in Predictive Maintenance	6
2.2.4	Physics-Informed Neural Networks	7
2.2.5	Physics-Informed Neural Networks for Bearing Predictive Maintenance	8
2.3	Research Gap	8
3	Dataset and Research Methodology	9
3.1	Introduction	9
3.2	Experimental Setup	10
3.3	Dataset Description	11
3.4	Feature Extraction	12
3.5	RUL Target Labeling	14
3.6	Preprocessing	14
3.7	PCA Health Indicator	15
3.8	Model Definitions	16
3.9	Experiment Design	16
3.10	Evaluation Metrics	17
4	Local Implementation and Experimental Setup	18
4.1	Local Project Structure	18
4.2	Reproducibility Commands	18
4.3	Software Environment	18
4.4	Data Validation	19
4.5	Feature Cache	19

4.6	Visual Reporting	19
5	Results and Discussion	20
5.1	Feature Behavior Over Bearing Life	20
5.2	Main Prediction Results	22
5.3	Experiment-Wise Comparison	25
5.4	True Versus Predicted RUL	28
5.5	Error Analysis	30
5.6	Discussion of the Physics-Informed Model	31
5.7	Why the LSTM Performed Well	32
5.8	Why the CNN Won Experiment 3	32
5.9	Practical Implications	32
6	Conclusion	33
6.1	Summary	33
6.2	Contributions	33
6.3	Limitations	34
6.4	Future Work	34
6.5	Final Conclusion	34
	References	35
A	Reproducibility Notes	38
A.1	Final Result Table	38
A.2	Commands Used for the Local Pipeline	38
A.3	Raw Data Location	38

A.4	LaTeX Build	39
-----	-----------------------	----

List of Figures

3.1	End-to-end local workflow used for IMS bearing RUL prediction.	9
3.2	IMS bearing test rig and sensor placement used for run-to-failure data collection, adapted from the IMS Bearing Data README and Qiu et al. [1].	10
3.3	Snapshot coverage for the four bearing runs used in the final experiments.	12
3.4	Bearing fault-frequency values used as markers for envelope spectral-energy extraction.	14
3.5	Smoothed PCA health indicator compared with normalized damage on held-out test runs.	15
3.6	Experiment split matrix showing how each bearing run is used.	17
5.1	RMS evolution over normalized bearing life.	20
5.2	Kurtosis evolution over normalized bearing life.	21
5.3	Crest factor evolution over normalized bearing life.	21
5.4	Combined fault-frequency envelope energy evolution over normalized bearing life.	22
5.5	Average RMSE by model across the three experiments.	24
5.6	Average MAE by model across the three experiments.	24
5.7	Average R^2 by model across the three experiments.	25
5.8	Model ranking by RMSE for each experiment. Lower rank is better. . . .	26
5.9	Lowest-RMSE model in each experiment.	26
5.10	RMSE by model and experiment.	27
5.11	MAE by model and experiment.	27

5.12	R^2 by model and experiment.	28
5.13	Experiment 1 true and predicted normalized RUL on the held-out test run.	28
5.14	Experiment 2 true and predicted normalized RUL on the held-out test run.	29
5.15	Experiment 3 true and predicted normalized RUL on the held-out test run.	29
5.16	Smoothed absolute prediction error over normalized test life.	30
5.17	Signed prediction error distribution. Positive error means the model over-estimated normalized RUL.	31

List of Tables

3.1	Bearing runs used in the final experiments.	11
3.2	Extracted features used by the neural models.	13
3.3	Bearing fault frequencies used for envelope spectral-energy extraction. . .	13
3.4	PCA health-indicator monotonicity.	15
3.5	Model definitions and roles in the experiment.	16
3.6	Train-validation-test split for each experiment.	17
5.1	Final prediction metrics for all models and experiments.	23
5.2	Average model performance across the three experiments.	23
5.3	Model ranking by RMSE.	25
A.1	Final prediction metrics for all models and experiments.	38

Nomenclature

Symbol or term	Meaning
ANN	Artificial neural network
BPFO	Ball pass frequency outer race
BPFI	Ball pass frequency inner race
BSF	Ball spin frequency
CNN	Convolutional neural network
DeepXDE	Python library used for physics-informed neural networks
FFT	Fast Fourier transform
FNN	Feed-forward neural network
FTF	Fundamental train frequency
HI	Health indicator
IMS	Intelligent Maintenance Systems
LSTM	Long short-term memory network
MAE	Mean absolute error
PCA	Principal component analysis
PHM	Prognostics and health management
PINN	Physics-informed neural network
REB	Rolling-element bearing
RMS	Root mean square
RMSE	Root mean squared error
RUL	Remaining useful life
R^2	Coefficient of determination

Chapter 1

Introduction

1.1 BACKGROUND AND MOTIVATION

Rolling-element bearings support shafts, carry loads, and reduce friction in motors, pumps, turbines, gearboxes, fans, compressors, and many other machines. Because they sit directly in the load path, a bearing defect rarely remains an isolated issue for long. A small defect on a raceway or rolling element can create repeated impacts, excite vibration, produce heat, damage nearby parts, and eventually force the machine out of service. In industrial maintenance, the cost of a bearing failure is often much larger than the cost of the bearing itself.

Traditional maintenance strategies do not fully solve the problem. Reactive maintenance waits for failure and accepts the risk of unplanned shutdown. Preventive maintenance replaces parts on a fixed schedule, but that schedule is usually conservative. It can remove healthy components too early or miss faults that develop faster than expected. Condition-based maintenance improves on both approaches by measuring the actual condition of the machine and using that evidence to support maintenance decisions. Prognostics and health management extends this idea by estimating not only whether a component is damaged, but how much useful life may remain before failure [2, 3].

Remaining useful life prediction is difficult for bearings because degradation is not a simple straight line. Vibration can remain low for a long portion of the test, then change rapidly near the end of life. The measured signal depends on fault type, speed, load, sensor placement, noise, resonance, lubrication, and bearing geometry [4, 5]. A model trained on one bearing may not behave well on another bearing, even when both bearings come from the same test rig. This is why benchmark run-to-failure datasets such as the IMS bearing dataset are useful. They allow researchers to study the same raw measurements and compare methods under repeatable conditions [1].

Deep learning methods have become common in bearing diagnosis and RUL prediction because they can model nonlinear patterns in vibration signals. CNNs can learn local patterns from windows of data, LSTMs can model temporal dependence, and feed-forward networks can fit engineered feature vectors [6–8]. These methods can work well when the

training and test data are similar. Their weakness is that they usually learn only from examples. When labeled run-to-failure data are limited or the test bearing degrades differently from the training bearings, a data-only model may produce predictions that are numerically plausible but physically questionable.

Physics-informed neural networks offer a way to include prior engineering knowledge in the learning process. In their most established form, PINNs train a neural network while penalizing violations of governing equations or physical constraints [9]. Bearing RUL prediction does not usually provide a complete closed-form degradation equation, but it does provide useful partial knowledge. Damage should generally increase with operating time, RUL should generally decrease, and bearing geometry gives characteristic fault frequencies associated with cage, inner-race, outer-race, and rolling-element defects. This thesis studies whether such weak physical information can improve RUL prediction on the IMS data.

1.2 PROBLEM STATEMENT

The problem studied in this thesis is the prediction of normalized remaining useful life for selected bearings in the IMS run-to-failure dataset. The main question is whether a physics-informed neural network using time and fault-frequency information can generalize better across bearing runs than comparable data-driven neural baselines.

This question is practical as well as methodological. In a real machine, the future failure path is unknown. A maintenance engineer may have historical data from one bearing, but the next bearing may not degrade in the same way. A useful RUL model must therefore do more than memorize one run. It must learn a relationship between measured condition and remaining life that transfers, at least partly, across bearings.

1.3 OBJECTIVES OF THE STUDY

The objectives of this thesis are:

1. To extract degradation-relevant vibration features from selected IMS bearing runs and construct normalized RUL targets with a PCA-based health indicator for analysis.
2. To design and evaluate a DeepXDE physics-informed neural model that uses weak monotonicity and fault-frequency energy priors for bearing RUL prediction.
3. To compare the proposed physics-informed model with data-only, LSTM, and CNN neural baselines under complete held-out bearing splits and interpret where each model succeeds or fails.

1.4 SCOPE

This study is limited to the IMS bearing data placed locally under `data/raw/1st_test`, `data/raw/2nd_test`, and `data/raw/3rd_test`. Four bearing runs are used in the final experiment design: `ds2_b1`, `ds1_b3`, `ds1_b4`, and `ds3_b3`. The work uses vibration signals only. It does not add temperature, oil debris, acoustic emission, or motor-current data.

The models predict normalized RUL from extracted features. The thesis does not claim field deployment readiness. The data come from a controlled test rig, and the operational conditions are much simpler than many industrial machines. The physical information in the PINN is also intentionally weak. It uses monotonicity and fault-frequency energy priors rather than a full contact mechanics or crack growth model.

1.5 IMPORTANCE OF THE STUDY

This work connects three parts of the bearing prognostics problem in one local workflow: signal processing, machine learning, and physical interpretation. The feature extraction step keeps the connection to vibration analysis through RMS, kurtosis, crest factor, and envelope spectral energies. The neural models test whether those features can support RUL prediction across bearing runs. The physics-informed model tests whether simple physical priors can regularize the learning process.

The final result is not a simple success story for the proposed model. That is important. A thesis result should report what the experiment shows, not what the proposal expected. The completed run suggests that LSTM sequence modeling captured the degradation trend more consistently than the proposed weakly physics-informed feed-forward model. At the same time, the PINN was competitive in Experiment 2 and ranked first there by RMSE. This mixed behavior gives a clearer direction for future work than an overconfident positive claim would.

1.6 LIMITATIONS

The main limitations are:

1. The IMS dataset contains a small number of run-to-failure records, so the experiment design is constrained.
2. The selected physics prior is not a full bearing degradation law. It is a weak regularizer based on monotonic RUL behavior and fault-frequency energy.
3. The model uses hand-engineered features, not raw vibration waveforms.

4. The RUL labels are derived from elapsed time to the end of each run, which assumes the last available file represents failure or near-failure.
5. The experiments are run on one local PC, so training time and numerical results may vary on other machines.

1.7 THESIS OUTLINE

Chapter 2 reviews bearing failure, vibration-based condition monitoring, RUL prediction, deep learning baselines, and physics-informed learning. Chapter 3 describes the IMS dataset, feature extraction, preprocessing, health indicator construction, model architectures, and evaluation metrics. Chapter 4 describes the local implementation, experiment splits, software environment, and reproducibility setup. Chapter 5 presents the results and discusses model behavior across the three experiments. Chapter 6 summarizes the findings, contributions, limitations, and future work.

Chapter 2

Literature Review

2.1 INTRODUCTION

This chapter reviews the research background for bearing RUL prediction with physics-informed neural networks. The discussion follows the structure of the thesis proposal, but it is updated to reflect the completed experiments. It begins with rolling-element bearing failure modes and vibration-based health monitoring, then reviews predictive maintenance methods, deep learning for bearing prognosis, general PINN theory, and recent physics-informed learning work in bearing condition monitoring. The chapter closes by identifying the research gap addressed in this thesis.

2.2 LITERATURE SURVEY

2.2.1 Rolling-Element Bearings: Fundamentals and Failure Modes

A rolling-element bearing contains an inner race, an outer race, rolling elements, and a cage. The rolling elements transfer load through small contact regions, so repeated stress cycles can initiate subsurface or surface-initiated fatigue. Defects may also arise from lubrication failure, contamination, overload, misalignment, electrical pitting, poor installation, or manufacturing variation. Once a defect begins, repeated rolling contact can grow it into a spall or crack that changes the vibration response of the machine [10–12].

The vibration signature depends on the component that is damaged. An outer-race defect produces impacts when rolling elements pass over a fixed damaged region. An inner-race defect rotates with the shaft and is affected by the load zone. Rolling-element and cage defects produce different modulation patterns. These mechanisms explain why bearing monitoring often uses characteristic frequencies derived from bearing geometry and shaft speed. The frequencies are not perfect fault labels, but they provide physically meaningful regions in the spectrum and envelope spectrum [4, 13].

The thesis proposal emphasized that bearing prognostics should not treat vibration features as arbitrary numbers. That point remains important in the completed work. RMS, kurtosis, crest factor, and envelope energy have mechanical meaning, even when they

are later passed into neural models. Using physically interpretable features also makes it easier to see why a model succeeds or fails on a held-out bearing run.

2.2.2 Traditional Predictive Maintenance Techniques for Bearings

Predictive maintenance is normally discussed as part of condition-based maintenance and prognostics and health management. Jardine et al. [2] describe the broader machinery diagnostics and prognostics framework, while Lei et al. [3] review the full path from data acquisition to RUL prediction. In a bearing application, the workflow usually includes sensor measurement, preprocessing, feature extraction, health indicator construction, fault diagnosis, and future-life estimation.

Vibration remains one of the most widely used signals for bearing health monitoring because bearing faults generate impacts and modulations in the mechanical response. Time-domain features such as RMS and kurtosis summarize energy and impulsiveness. Frequency-domain and envelope-domain features can reveal periodic impact patterns linked to bearing kinematics. Similarity-based and health-indicator methods then convert these features into a degradation measure or RUL estimate [14].

The IMS bearing dataset is one of the common public run-to-failure datasets used for this type of work. It is associated with the wavelet-filter bearing prognostics study of Qiu et al. [1]. The PRONOSTIA platform is another well-known accelerated bearing degradation benchmark [15]. Public datasets are valuable because they let different researchers test methods on the same raw measurements. At the same time, public datasets are small compared with industrial variability, so the train-test split must be designed carefully. Random snapshot splits can overstate performance by allowing information from the same run-to-failure trajectory into both training and testing.

2.2.3 Deep Learning in Predictive Maintenance

Deep learning has become popular in bearing prognostics because it can fit nonlinear relationships between sensor measurements and degradation labels. A feed-forward neural network can work directly with engineered feature vectors. A CNN can detect local patterns from a sequence or transformed signal representation. An LSTM can model temporal dependence and is therefore naturally suited to degradation trajectories. Babu et al. [6] showed early CNN-based RUL regression, and later studies extended deep learning for transfer learning, high-quality representation learning, and hybrid sequence architectures [7, 8, 16].

Recent studies also combine convolutional, recurrent, and autoencoder structures for

bearing health monitoring and RUL prediction [17–19]. These models are attractive because inference is fast after training and because they can learn patterns that are difficult to specify manually. However, their performance depends strongly on the similarity between the training and test distributions. In bearing RUL prediction, this is a serious issue because two bearings under similar nominal conditions may fail at different times and with different feature trajectories.

For this reason, model evaluation should hold out complete bearing runs when the research question is transfer across bearings. A random split of individual snapshots may measure interpolation within one degradation curve rather than generalization to another bearing. The experiment design in this thesis uses complete held-out runs for validation and testing.

2.2.4 Physics-Informed Neural Networks

Physics-informed neural networks were introduced as a way to train neural networks under physical residual constraints [9]. In the usual PINN setting, a neural network approximates a field variable, and automatic differentiation is used to compute derivatives that appear in a governing equation. The training loss combines data mismatch and a physics residual. This has made PINNs attractive in fluid mechanics, heat transfer, wave propagation, and inverse problems where governing equations are available.

The PINN literature has expanded quickly. Reviews by Lawal et al. [20] and Ren et al. [21] describe methodological growth, optimization issues, and applications across scientific computing. The strength of a PINN is not simply that it contains a neural network and a physical equation. Its value depends on whether the physical residual is relevant, correctly weighted, and numerically trainable. Poorly chosen physical constraints can slow training or add bias.

Bearing RUL prediction is not a textbook PINN problem because the degradation process is not governed by a single known partial differential equation. Rolling contact fatigue, lubrication, load distribution, resonance, material variation, and measurement noise interact in ways that are difficult to reduce to one residual. Crack-growth laws such as Paris and Erdogan [22] show how mechanics can describe damage propagation in some contexts, but a complete bearing RUL equation for vibration features is rarely available in practice. This motivates weaker physics-informed learning: monotonic damage behavior, lifetime distributions, fault-frequency consistency, and other engineering priors.

2.2.5 Physics-Informed Neural Networks for Bearing Predictive Maintenance

Recent bearing-related studies have begun to use physics-informed or knowledge-informed learning. Chen et al. [23] proposed a physics-informed deep neural network for bearing prognosis with multisensory signals. Parziale et al. [24] applied PINNs to condition monitoring of rotating shafts. Herwig et al. [25] developed a signal-processing and physics-informed network for explainable bearing condition monitoring. Zhong et al. [26] used a multi-source physics-informed SincNet structure for rolling bearing fault diagnosis. von Hahn and Mechefske [27] investigated a knowledge-informed loss based on Weibull lifetime behavior.

These studies show why the approach is promising. Physics can improve interpretability, guide learning when data are limited, and discourage predictions that conflict with engineering expectations. They also show why the approach must be tested carefully. A physical prior only helps when it matches the measured data closely enough. If the prior is too weak, the model may behave like a data-only network. If it is too strong or mismatched, it may prevent the model from fitting valid degradation patterns.

2.3 RESEARCH GAP

The literature shows three gaps that motivate this thesis. First, many bearing RUL studies report strong results, but their preprocessing, RUL labeling, train-test split, and evaluation protocol differ. This makes comparisons difficult. Second, some studies use random sample splits that can leak run-specific degradation information into both training and testing. Third, physics-informed learning for bearing prognosis is still developing, and there is no guarantee that a weak physical prior improves prediction on complete held-out bearing runs.

This thesis addresses a narrower and reproducible question: how does a weak physics-informed RUL model behave when compared with common neural baselines on complete held-out IMS bearing runs? The study uses a local project structure, explicit feature extraction, a defined RUL target, saved tables and figures, and three train-validation-test splits. The goal is not to claim a universal best model. The goal is to test a specific physics-informed formulation and report what the completed experiment shows.

Chapter 3

Dataset and Research Methodology

3.1 INTRODUCTION

This chapter describes the dataset, signal-processing steps, RUL target construction, model definitions, experiment design, and evaluation metrics used in the study. The workflow begins with raw IMS vibration files and ends with model rankings, prediction curves, and error analysis. Figure 3.1 summarizes the local pipeline.

Methodology workflow

Local project stages used to generate the thesis results.




Figure 3.1: End-to-end local workflow used for IMS bearing RUL prediction.

3.2 EXPERIMENTAL SETUP

The IMS bearing data were generated by the NSF I/UCR Center for Intelligent Maintenance Systems with support from Rexnord Corporation. The test rig used four Rexnord ZA-2115 double-row bearings mounted on a common shaft. An AC motor kept the shaft speed constant at 2000 RPM, and a spring loading mechanism applied a radial load of 6000 lb to the shaft and bearings. The bearings were force lubricated throughout the run-to-failure tests. The IMS README also notes that the observed failures occurred after the bearings exceeded their designed lifetime of more than 100 million revolutions.

Vibration was measured at the bearing housings using PCB 353B33 High Sensitivity Quartz ICP accelerometers. The first data set used two accelerometers per bearing, placed in the x- and y-axis directions. The second and third data sets used one accelerometer per bearing. Data collection was performed with an NI DAQ Card 6062E. Figure 3.2 shows the bearing test rig and sensor placement used for the IMS experiments.




Figure 3.2: IMS bearing test rig and sensor placement used for run-to-failure data collection, adapted from the IMS Bearing Data README and Qiu et al. [1].

3.3 DATASET DESCRIPTION

The IMS bearing dataset contains run-to-failure vibration measurements collected from the setup described above. The raw data are stored as timestamped ASCII text files. Each file contains a nominal one-second vibration snapshot with 20,480 data points sampled at 20 kHz. The filename records the collection time, and larger gaps between timestamps indicate that the experiment resumed on the next working day. In this project, the extracted dataset folders are stored locally under `data/raw`.

The local pipeline uses three IMS folders: `1st_test`, `2nd_test`, and `3rd_test`. The first test contains eight columns because two accelerometer channels are available for each of four bearings. The second and third tests contain four columns, one for each bearing. The implementation maps each selected column to a bearing and averages two axes when the first test provides x and y directions for the target bearing.

Table 3.1: Bearing runs used in the final experiments.

Run ID	IMS dataset	Bearing number	Number of samples	Used in experiments
ds2_b1	2nd_test	1	984	Exp 1 Train; Exp 2 Train; Exp 3 Train
ds1_b3	1st_test	3	2156	Exp 1 Train; Exp 2 Validation; Exp 3 Test
ds1_b4	1st_test	4	2156	Exp 1 Validation; Exp 2 Train; Exp 3 Validation
ds3_b3	3rd_test	3	6324	Exp 1 Test; Exp 2 Test; Exp 3 Train

Figure 3.3 shows the number of snapshots available for each selected run. The difference in run length is large. The third dataset bearing run has 6324 samples, while the second dataset bearing run has 984 samples. This imbalance affects model training and is one reason why Experiment 3 uses balanced sampling for its training runs.

Snapshot coverage by bearing run
 Counts reflect the extracted IMS folders used in the local pipeline.




Figure 3.3: Snapshot coverage for the four bearing runs used in the final experiments.

3.4 FEATURE EXTRACTION

Each vibration snapshot is converted into a feature vector. The time-domain features are RMS, standard deviation, peak-to-peak value, kurtosis, crest factor, and mean absolute value. These features summarize amplitude, spread, impulsiveness, and peak behavior. They are widely used in bearing condition monitoring because they are simple and physically interpretable.

The physics-related features are calculated from envelope spectral energy near bearing fault frequencies. The signal is centered, transformed with the Hilbert transform to obtain the envelope, centered again, windowed with a Hann window, and transformed with a real FFT. For each fault frequency, the pipeline sums spectral energy in a 5 Hz band around the first four harmonics. The four energy features are E_{FTF} , E_{BPFO} , E_{BPFI} , and E_{BSF} . Their sum is stored as E_{kin} .

The fault-frequency equations used by the implementation are

$$FTF = \frac{f_r}{2} \left(1 - \frac{d}{D} \cos \theta \right), \quad (3.1)$$

$$BPFO = \frac{nf_r}{2} \left(1 - \frac{d}{D} \cos \theta \right), \quad (3.2)$$

$$BPFI = \frac{nf_r}{2} \left(1 + \frac{d}{D} \cos \theta \right), \quad (3.3)$$

$$BSF = \frac{Df_r}{2d} \left(1 - \left(\frac{d}{D} \cos \theta \right)^2 \right), \quad (3.4)$$

where f_r is shaft frequency, n is the number of rolling elements, d is rolling-element diameter, D is pitch diameter, and θ is contact angle.

Table 3.2: Extracted features used by the neural models.

Feature name	Domain	Physical meaning	Used by models
elapsed_scaled	time feature	Normalized elapsed operating time used as temporal input	yes
rms	time-domain	Signal energy amplitude through root mean square vibration	yes
std	time-domain	Vibration dispersion around the mean	yes
ptp	time-domain	Peak-to-peak vibration range	yes
kurtosis	time-domain	Impulsiveness of vibration signal	yes
crest_factor	time-domain	Peak amplitude relative to RMS	yes
mean_abs	time-domain	Mean absolute vibration amplitude	yes
E_FTF	physics-informed frequency feature	Envelope spectral energy around cage fault train frequency harmonics	yes
E_BPFO	physics-informed frequency feature	Envelope spectral energy around outer-race fault frequency harmonics	yes
E_BPFI	physics-informed frequency feature	Envelope spectral energy around inner-race fault frequency harmonics	yes
E_BSF	physics-informed frequency feature	Envelope spectral energy around rolling-element fault frequency harmonics	yes
E_kin	physics-informed frequency feature	Combined bearing kinematic fault-frequency spectral energy	yes
sigma_H_norm	auxiliary placeholder	Constant placeholder; not a Hertzian contact stress calculation	yes

Table 3.3: Bearing fault frequencies used for envelope spectral-energy extraction.

Frequency name	Symbol	Value in Hz	Related bearing component	Extracted feature
FTF	FTF	14.78	Cage / fundamental train	E_FTF
BPFO	BPFO	236.40	Outer race	E_BPFO
BPFI	BPFI	296.93	Inner race	E_BPFI
BSF	BSF	139.92	Rolling element	E_BSF

Bearing fault-frequency markers
Kinematic frequencies used for envelope spectral-energy features.




Figure 3.4: Bearing fault-frequency values used as markers for envelope spectral-energy extraction.

3.5 RUL TARGET LABELING

For each selected bearing run, elapsed time is computed from the timestamp of the first snapshot. RUL is computed by subtracting the current timestamp from the timestamp of the final snapshot. The model target is normalized RUL:

$$RUL_{norm}(t) = \frac{RUL(t)}{\max(RUL)}. \quad (3.5)$$

This gives a target near 1 at the beginning of the run and near 0 at the end. The normalization makes results comparable across runs with different durations. It also means that the model predicts relative life fraction, not absolute hours.

3.6 PREPROCESSING

Feature values can differ greatly between bearing runs. To reduce the effect of run-specific scale, each run is normalized using its early healthy baseline. The first 5 percent of samples, with a minimum of 20 samples, are treated as the healthy baseline. For each model feature, the median baseline value is used. The feature is converted to a nonnegative relative increase, transformed with $\log_{1p}$, and smoothed with a rolling median of 15 samples.

Elapsed time is converted to `elapsed_scaled` by dividing elapsed hours by 1200. The model input vector contains `elapsed_scaled`, the normalized feature set, and

`sigma_H_norm`. The final variable is a constant placeholder kept for compatibility with the original notebook structure. It is not treated as a real Hertzian contact stress calculation in the final interpretation.

3.7 PCA HEALTH INDICATOR

A PCA-based health indicator is built to inspect whether the extracted features contain a monotonic degradation signal. The model feature matrix is standardized and projected onto the first principal component. The resulting component is scaled between 0 and 1. If the indicator direction is inverted, it is flipped so that higher values correspond to greater damage.

The health indicator is not used as the final RUL target. It is used as an explanatory tool. It helps show whether the engineered features broadly track degradation.

Table 3.4: PCA health-indicator monotonicity.

Bearing run	Monotonicity score	Interpretation
ds1_b3	0.766	Strongly degradation-consistent PCA-HI behavior
ds3_b3	0.864	Strongly degradation-consistent PCA-HI behavior




Figure 3.5: Smoothed PCA health indicator compared with normalized damage on held-out test runs.

3.8 MODEL DEFINITIONS

Four models are evaluated. The data-only neural baseline is a feed-forward neural network implemented in DeepXDE [28]. It uses three hidden layers with 64 neurons per layer and tanh activation. A sigmoid output transform constrains predictions between 0 and 1. It is trained with mean squared error.

The proposed physics-informed model uses the same basic feed-forward architecture but adds weak physical residuals. The first residual penalizes positive RUL slope with respect to time beyond a small tolerance, because RUL should generally decrease as life is consumed. The second residual connects high E_{kin} to damage by encouraging the predicted damage term $1 - RUL$ to be consistent with fault-frequency energy. The model loss can be written as

$$\mathcal{L} = \mathcal{L}_{data} + \lambda_m \mathcal{L}_{mono} + \lambda_e \mathcal{L}_{energy}, \quad (3.6)$$

where $\mathcal{L}_{data}$ is the supervised RUL loss, $\mathcal{L}_{mono}$ is the monotonicity residual, $\mathcal{L}_{energy}$ is the fault-energy residual, and λ_m and λ_e are small regularization weights.

The LSTM baseline uses sliding windows of feature vectors. It reads a sequence of length 20 and predicts the RUL at the end of the sequence. The architecture has one LSTM layer with 64 hidden units followed by a small dense head and a sigmoid output. The CNN baseline uses the same sequence windows but applies 1D convolutions across the time axis.

Table 3.5: Model definitions and roles in the experiment.

Model	Model type	Input format	Role in thesis	Target
Data-only neural baseline	Feed-forward neural network	Single snapshot feature vector	Data-only baseline	Normalized RUL
Proposed DeepXDE Physics-Informed RUL Model	DeepXDE physics-informed neural network	Single snapshot feature vector with time and fault-frequency energy features	Proposed model	Normalized RUL
LSTM baseline	Recurrent neural network	Sliding-window sequence of feature vectors	Sequence baseline	Normalized RUL
CNN baseline	1D convolutional neural network	Sliding-window sequence of feature vectors	Sequence baseline	Normalized RUL

3.9 EXPERIMENT DESIGN

The three experiments hold out complete bearing runs for validation and testing. This is stricter than randomly splitting snapshots because the model must transfer across bearing runs. Table 3.6 and Figure 3.6 show the split design.

Table 3.6: Train-validation-test split for each experiment.

Experiment	Train runs	Validation run	Test run	Purpose
Exp 1	ds2_b1 + ds1_b3	ds1_b4	ds3_b3	Cross-dataset/cross-bearing transfer with ds3_b3 held out for testing
Exp 2	ds2_b1 + ds1_b4	ds1_b3	ds3_b3	Alternative training-bearing combination with the same ds3_b3 test run
Exp 3	ds2_b1 + ds3_b3	ds1_b4	ds1_b3	Cross-dataset transfer to ds1_b3 with balanced training runs




Figure 3.6: Experiment split matrix showing how each bearing run is used.

3.10 EVALUATION METRICS

The models are evaluated with MAE, RMSE, and R^2 . MAE is the average absolute prediction error. RMSE penalizes larger errors more strongly because the error is squared before averaging. R^2 measures how much variance in the target is explained by the prediction. A negative R^2 means the model performs worse than predicting the mean target value.

The thesis uses RMSE as the main ranking metric because large RUL errors are especially undesirable in maintenance planning. MAE and R^2 are still reported to give a fuller view of model behavior. The metric definitions follow common prognostics evaluation practice [29].

Chapter 4

Local Implementation and Experimental Setup

4.1 LOCAL PROJECT STRUCTURE

The original work began as a Google Colab notebook. For this thesis, the workflow was converted into a local Python project. The code no longer mounts Google Drive. It reads data from `data/raw`, caches extracted features in `data/processed_features`, and writes tables and figures to `outputs`.

The main package is under `src/thesis_work`. It contains separate modules for configuration, IMS data loading, feature extraction, preprocessing, model training, metrics, report generation, and command-line execution. This structure makes the experiment easier to test and rerun than a single notebook.

4.2 REPRODUCIBILITY COMMANDS

The project uses `uv` for dependency management. The main commands are:

```
uv sync --extra dev
uv run thesis-work validate-data
uv run thesis-work extract-features
uv run thesis-work run
uv run thesis-work regenerate-figures
uv run pytest -q
```

The full run trains all four models for all three experiments. The command can take a long time on CPU, so the project also supports shorter smoke-test arguments. The final results in this thesis come from the completed full run saved in `outputs/tables/final_results_table.csv`.

4.3 SOFTWARE ENVIRONMENT

The implementation uses Python 3.11 and the following main libraries: NumPy, pandas, SciPy, scikit-learn, matplotlib, seaborn, DeepXDE, PyTorch, and tqdm. The LaTeX

thesis is compiled with MiKTeX using the Windows batch script included in the thesis source folder.

The local PC used for the completed run had an AMD Ryzen 7 5800H CPU with 8 cores and 16 logical processors, 16 GB RAM, Windows 11 Home Single Language, and an NVIDIA GeForce RTX 3060 Laptop GPU with 4 GB VRAM. Training and inference wall time is machine-specific and should be reported together with this hardware context.

4.4 DATA VALIDATION

The validation command checks that the expected folders are present and that the first snapshot in each dataset can be read with the correct number of columns. This step matters because the IMS text files may be stored either as extracted folders or zip files. The local loader supports both, but the final run used extracted folders under `data/raw`.

4.5 FEATURE CACHE

Feature extraction is the slowest preprocessing step because every vibration snapshot must be read and transformed. The pipeline therefore caches per-run feature tables and a combined `all_runs_features.csv`. Cached features make it possible to regenerate figures and tables without reprocessing the raw data or retraining the models.

4.6 VISUAL REPORTING

The final figures were generated with matplotlib and seaborn using a consistent color palette. Models use the same colors across all plots: gray for the data-only baseline, blue for the proposed PINN, green for LSTM, and orange-red for CNN. Bearing runs also use fixed colors. In the LaTeX thesis, plot files avoid embedded chart titles where possible because figure captions and labels provide the formal titles in the document.

Chapter 5

Results and Discussion

5.1 FEATURE BEHAVIOR OVER BEARING LIFE

The extracted features do not evolve identically across bearing runs. This is expected because different bearings can have different fault initiation times and degradation rates. RMS and peak-to-peak values capture energy growth. Kurtosis and crest factor are more sensitive to impulsive behavior. E_{kin} summarizes energy near the calculated fault-frequency harmonics.

Figures 5.1 to 5.4 show representative feature evolution over normalized lifetime. The plotted feature values are normalized for visualization so that their trends can be compared across bearings.




Figure 5.1: RMS evolution over normalized bearing life.

Kurtosis evolution over bearing life

Features are normalized by each run's early healthy baseline.




Figure 5.2: Kurtosis evolution over normalized bearing life.

Crest factor evolution over bearing life

Features are normalized by each run's early healthy baseline.




Figure 5.3: Crest factor evolution over normalized bearing life.

Combined fault-frequency energy E_{kin} evolution over bearing life

Features are normalized by each run's early healthy baseline.



Figure 5.4: Combined fault-frequency envelope energy evolution over normalized bearing life.

The plots show why cross-bearing RUL prediction is difficult. Some features rise smoothly in one run and irregularly in another. A model trained on one degradation shape may misread a different shape. This is the central challenge behind the results that follow.

5.2 MAIN PREDICTION RESULTS

Table 5.1 gives the full model results for all three experiments.

Table 5.1: Final prediction metrics for all models and experiments.

Experiment	Model	MAE	RMSE	R2
Exp 1	Data-only FNN	0.083	0.128	0.801
Exp 1	Proposed PINN	0.147	0.173	0.638
Exp 1	LSTM	0.101	0.112	0.848
Exp 1	CNN	0.244	0.266	0.146
Exp 2	Data-only FNN	0.104	0.148	0.738
Exp 2	Proposed PINN	0.114	0.144	0.751
Exp 2	LSTM	0.141	0.155	0.708
Exp 2	CNN	0.200	0.235	0.334
Exp 3	Data-only FNN	0.156	0.280	0.032
Exp 3	Proposed PINN	0.262	0.351	-0.530
Exp 3	LSTM	0.176	0.219	0.382
Exp 3	CNN	0.183	0.204	0.463

The average performance across the three experiments is shown in Table 5.2. On average, the LSTM baseline is the strongest model. It has the lowest average RMSE at 0.162 and the highest average R^2 at 0.646. The data-only feed-forward baseline is second by average RMSE at 0.185. The proposed physics-informed model has an average RMSE of 0.223. The CNN baseline has the weakest average RMSE at 0.235, although it performs best in Experiment 3.

Table 5.2: Average model performance across the three experiments.

Model	Average MAE	Average RMSE	Average R2
CNN	0.209	0.235	0.314
Data-only FNN	0.114	0.185	0.524
LSTM	0.139	0.162	0.646
Proposed PINN	0.175	0.223	0.286




Figure 5.5: Average RMSE by model across the three experiments.




Figure 5.6: Average MAE by model across the three experiments.




Figure 5.7: Average R^2 by model across the three experiments.

These results do not support the original expectation that the proposed PINN would dominate the baselines. They show a more specific conclusion: the weak physics-informed formulation can help in one split, but it is not robustly better across all held-out bearing runs.

5.3 EXPERIMENT-WISE COMPARISON

The model ranking by RMSE is given in Table 5.3 and Figure 5.8. The ranking figure is oriented so that ranks increase upward: rank 1 is at the bottom and rank 4 is at the top.

Table 5.3: Model ranking by RMSE.

Experiment	Model	RMSE	Rank by RMSE
Exp 1	LSTM	0.112	1
Exp 1	Data-only FNN	0.128	2
Exp 1	Proposed PINN	0.173	3
Exp 1	CNN	0.266	4
Exp 2	Proposed PINN	0.144	1
Exp 2	Data-only FNN	0.148	2
Exp 2	LSTM	0.155	3
Exp 2	CNN	0.235	4
Exp 3	CNN	0.204	1
Exp 3	LSTM	0.219	2
Exp 3	Data-only FNN	0.280	3
Exp 3	Proposed PINN	0.351	4




Figure 5.8: Model ranking by RMSE for each experiment. Lower rank is better.

The ranking changes from experiment to experiment. In Experiment 1, the LSTM baseline performs best with RMSE 0.112. In Experiment 2, the proposed PINN performs best with RMSE 0.144, narrowly ahead of the data-only baseline. In Experiment 3, the CNN baseline performs best with RMSE 0.204, followed by the LSTM. The proposed PINN performs worst in Experiment 3, with negative R^2 .

Figure 5.9 makes this instability clear by showing only the lowest-RMSE model in each experiment. A single model does not win every split. The metric comparison plots in Figures 5.10 to 5.12 show the same pattern from different viewpoints.




Figure 5.9: Lowest-RMSE model in each experiment.

RMSE by model and experiment




Figure 5.10: RMSE by model and experiment.

MAE by model and experiment




Figure 5.11: MAE by model and experiment.




Figure 5.12: R^2 by model and experiment.

5.4 TRUE VERSUS PREDICTED RUL

Figures 5.13 to 5.15 compare the true normalized RUL curve with model predictions on the held-out test runs.




Figure 5.13: Experiment 1 true and predicted normalized RUL on the held-out test run.

Exp 2: true versus predicted RUL




Figure 5.14: Experiment 2 true and predicted normalized RUL on the held-out test run.

Exp 3: true versus predicted RUL




Figure 5.15: Experiment 3 true and predicted normalized RUL on the held-out test run.

In Experiment 1, the LSTM follows the decreasing RUL trend more closely than the other models. The proposed PINN is smoother, but it does not match the test curve as well. In Experiment 2, the PINN is more competitive and gives the lowest RMSE. This suggests that the physics prior can help when the training bearings give a degradation pattern compatible with the held-out test bearing. In Experiment 3, the PINN struggles on `ds1_b3`, while the CNN and LSTM sequence models handle the test run better.

5.5 ERROR ANALYSIS

Figure 5.16 shows smoothed absolute prediction error over normalized test life. The largest errors tend to appear when the degradation curve changes more sharply or when a model predicts RUL too conservatively near the end of life.

Smoothed absolute prediction error over test life

Errors are aligned to each model's prediction positions.




Figure 5.16: Smoothed absolute prediction error over normalized test life.




Figure 5.17: Signed prediction error distribution. Positive error means the model overestimated normalized RUL.

The distribution confirms that the sequence models do not merely improve the average score. They also reduce the spread of errors in several cases. The proposed PINN has a wider error distribution because its weak physical prior does not fully capture the different degradation paths of the test bearings.

5.6 DISCUSSION OF THE PHYSICS-INFORMED MODEL

The proposed PINN includes two forms of physical regularization: a monotonicity residual and a spectral prior based on fault-frequency energy. Both are reasonable engineering assumptions, but both are incomplete.

The monotonicity residual assumes RUL should not increase with elapsed life. That is correct at the life-label level, but the model sees noisy features. If the feature pattern of a held-out bearing differs from the training bearings, a strict preference for smooth monotonic behavior can prevent the model from adapting to local signal changes. The implementation uses a tolerance and small weight to reduce this issue, but the result still depends on the split.

The spectral prior assumes that higher combined fault-frequency energy should correspond to greater damage. This is physically plausible, but envelope energy is not a perfect damage variable. It can be affected by resonance, noise, axis direction, load-zone behavior, and the timing of fault development. If E_{kin} grows late or irregularly in one run, a model trained on another run may apply the prior at the wrong time.

The strongest lesson is that physics-informed learning is not automatically better because it contains engineering language. The prior must match the data-generating process closely enough to help. In this thesis, the prior is useful in Experiment 2 but not enough to beat sequence modeling overall.

5.7 WHY THE LSTM PERFORMED WELL

The LSTM baseline uses recent history instead of a single snapshot. This is important because degradation is a process. A feature value by itself may be ambiguous, but its recent trend gives more information. For example, moderate RMS may mean early healthy operation in one run or a temporary plateau after fault growth in another. A short sequence helps the model distinguish these cases.

The LSTM also has fewer assumptions about the exact physical relationship between E_{kin} and damage. It learns temporal patterns directly from the feature sequences. This flexibility likely explains why it achieved the best average RMSE and R^2 .

5.8 WHY THE CNN WON EXPERIMENT 3

The CNN baseline performs worst on average but best in Experiment 3. This result shows that model choice depends on the train-test pairing. A 1D CNN can detect local temporal patterns in the sliding feature window. In Experiment 3, those local patterns transferred well to `ds1_b3`. In the other experiments, the same model did not generalize as strongly.

This does not mean the CNN is the best overall choice. It means that a single average metric can hide split-specific behavior. For a thesis, the split-wise analysis is therefore as important as the average table.

5.9 PRACTICAL IMPLICATIONS

For an engineering maintenance system, the safest conclusion is conservative. The current PINN formulation should not be deployed as a final RUL predictor. The LSTM baseline is stronger in the completed experiments, but it also needs more validation before field use. A practical system would need more run-to-failure data, multiple operating conditions, uncertainty estimates, and a clear maintenance decision rule.

The work still gives useful direction. Physics-informed features are valuable for interpretation, and the PINN framework remains promising. The next step is not to abandon physics. It is to improve the physical model, use better degradation constraints, and evaluate on more splits and datasets.

Chapter 6

Conclusion

6.1 SUMMARY

This thesis studied remaining useful life prediction for rolling-element bearings using the IMS run-to-failure dataset. A local Python project was built from the original notebook so the workflow can run without Google Drive. The pipeline reads raw IMS folders, extracts time-domain and envelope spectral features, constructs normalized RUL labels, builds a PCA health indicator, trains four neural models, and writes reproducible tables and figures.

The proposed model was a DeepXDE physics-informed neural network with weak physical residuals based on monotonic RUL behavior and fault-frequency energy. It was compared with a data-only feed-forward neural baseline, an LSTM sequence baseline, and a CNN sequence baseline across three cross-bearing experiments.

The results were mixed. The proposed PINN achieved the best RMSE in Experiment 2, but it did not perform best overall. The LSTM baseline had the best average performance, with average RMSE 0.162 and average R^2 0.646. The data-only baseline was second by average RMSE. The proposed PINN ranked third by average RMSE, and the CNN ranked fourth, although the CNN performed best in Experiment 3.

6.2 CONTRIBUTIONS

The main contributions of this thesis are:

1. A complete local RUL prediction pipeline for the IMS bearing dataset.
2. A documented feature extraction method combining time-domain vibration features and bearing fault-frequency envelope energy.
3. A PCA-based health indicator analysis for degradation consistency.
4. A DeepXDE physics-informed RUL model using monotonicity and spectral priors.

5. A fair comparison with feed-forward, LSTM, and CNN baselines under complete held-out bearing splits.
6. A clear discussion of the mismatch between the original proposal expectation and the final experimental evidence.

6.3 LIMITATIONS

The study has several limitations. The IMS data provide only a small number of complete failure runs. The RUL target is normalized by the final timestamp, so it does not represent an independently measured failure threshold. The physical prior is weak and does not include a full contact fatigue model, lubrication model, temperature effect, or uncertainty estimate. The model is trained on engineered features rather than raw waveforms. Finally, the results are from one local implementation and one set of hyperparameters.

6.4 FUTURE WORK

Future work should test stronger physics-informed formulations. Possible directions include a learned degradation state constrained by monotonic damage growth, uncertainty-aware RUL prediction, Weibull or crack-growth inspired lifetime priors, and multi-task learning that predicts both health indicator and RUL. The experiments should also be repeated on PRONOSTIA and other bearing datasets so that conclusions do not depend on one benchmark.

The feature pipeline can also be improved. Time-frequency representations, adaptive envelope bands, bearing-specific resonance selection, and raw waveform sequence models may capture information lost by the current feature set. More careful hyperparameter tuning and repeated random seeds would make the comparison more statistically reliable.

6.5 FINAL CONCLUSION

The completed experiments show that physics-informed bearing RUL prediction is promising but not automatic. In this thesis, the proposed PINN improved the result in one experiment but did not beat the LSTM baseline overall. The best model was the one that used temporal context most effectively. A practical bearing prognostics system should therefore combine physical interpretation with sequence-aware learning, and it should be tested under complete run-level splits before strong claims are made.

References

- [1] H. Qiu, J. Lee, J. Lin, and G. Yu, “Wavelet filter-based weak signature detection method and its application on rolling element bearing prognostics,” *Journal of Sound and Vibration*, vol. 289, no. 4–5, pp. 1066–1090, 2006.
- [2] A. K. S. Jardine, D. Lin, and D. Banjevic, “A review on machinery diagnostics and prognostics implementing condition-based maintenance,” *Mechanical Systems and Signal Processing*, vol. 20, no. 7, pp. 1483–1510, 2006.
- [3] Y. Lei, N. Li, L. Guo, N. Li, T. Yan, and J. Lin, “Machinery health prognostics: A systematic review from data acquisition to RUL prediction,” *Mechanical Systems and Signal Processing*, vol. 104, pp. 799–834, 2018.
- [4] G. Wu, T. Yan, G. Yang, H. Chai, and C. Cao, “A review on rolling bearing fault signal detection methods based on different sensors,” *Sensors*, vol. 22, no. 21, p. 8330, 2022.
- [5] V. Kannan, T. Zhang, and H. Li, “A review of the intelligent condition monitoring of rolling element bearings,” *Machines*, vol. 12, no. 7, p. 484, 2024.
- [6] G. Sateesh Babu, P. Zhao, and X.-L. Li, “Deep convolutional neural network based regression approach for estimation of remaining useful life,” in *Database Systems for Advanced Applications*, 2016, pp. 214–228.
- [7] T. Berghout, L.-H. Mouss, T. Bentrchia, and M. Benbouzid, “A semi-supervised deep transfer learning approach for rolling-element bearing remaining useful life prediction,” *IEEE Transactions on Energy Conversion*, vol. 37, no. 2, pp. 1200–1210, 2022.
- [8] A. Ayman, A. Onsy, O. Attallah, H. Brooks, and I. Morsi, “Feature learning for bearing prognostics: A comprehensive review of machine/deep learning methods, challenges, and opportunities,” *Measurement*, vol. 245, p. 116589, 2025.
- [9] M. Raissi, P. Perdikaris, and G. E. Karniadakis, “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving non-linear partial differential equations,” *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019.
- [10] P. Rycerz, A. Olver, and A. Kadiric, “Propagation of surface initiated rolling contact fatigue cracks in bearing steel,” *International Journal of Fatigue*, vol. 97, pp. 29–38, 2017.

- [11] F. Sadeghi, B. Jalalahmadi, T. S. Slack, N. Rajee, and N. K. Arakere, “A review of rolling contact fatigue,” *Journal of Tribology*, vol. 131, no. 4, 2009.
- [12] F. Xu, N. Ding, N. Li, L. Liu, N. Hou, N. Xu, W. Guo, L. Tian, H. Xu, C.-M. L. Wu, X. Wu, and X. Chen, “A review of bearing failure modes, mechanisms and causes,” *Engineering Failure Analysis*, vol. 152, p. 107518, 2023.
- [13] Y. Li, C. Zhang, T. R. Kurfess, S. Danyluk, and S. Y. Liang, “Diagnostics and prognostics of a single surface defect on roller bearings,” *Proceedings of the Institution of Mechanical Engineers, Part C: Journal of Mechanical Engineering Science*, vol. 214, no. 9, pp. 1173–1185, 2000.
- [14] W. Huang, A. Farahat, and C. Gupta, “Similarity-based feature extraction from vibration data for prognostics,” *Annual Conference of the PHM Society*, vol. 12, no. 1, 2020.
- [15] P. Nectoux, R. Gouriveau, K. Medjaher, E. Ramasso, B. Chebel-Morello, N. Zerhouni, and C. Varnier, “PRONOSTIA: An experimental platform for bearings accelerated degradation tests,” in *IEEE International Conference on Prognostics and Health Management*, 2012, pp. 1–8. [Online]. Available: <https://hal.science/hal-00719503>
- [16] C. Wang, W. Jiang, L. Shi, and L. Zhang, “Rolling bearing remaining useful life prediction using deep learning based on high-quality representation,” *Scientific Reports*, vol. 15, no. 1, 2025.
- [17] U. Farooq, M. Ademola, and A. Shaalan, “Comparative analysis of machine learning models for predictive maintenance of ball bearing systems,” *Electronics*, vol. 13, no. 2, p. 438, 2024.
- [18] K. Han, W. Wang, and J. Guo, “Research on a bearing fault diagnosis method based on a CNN-LSTM-GRU model,” *Machines*, vol. 12, no. 12, p. 927, 2024.
- [19] L. Yang, Y. Jiang, K. Zeng, and T. Peng, “Rolling bearing remaining useful life prediction based on CNN-VAE-MBiLSTM,” *Sensors*, vol. 24, no. 10, p. 2992, 2024.
- [20] Z. K. Lawal, H. Yassin, D. T. C. Lai, and A. Che Idris, “Physics-informed neural network (PINN) evolution and beyond: A systematic literature review and bibliometric analysis,” *Big Data and Cognitive Computing*, vol. 6, no. 4, p. 140, 2022.
- [21] Z. Ren, S. Zhou, D. Liu, and Q. Liu, “Physics-informed neural networks: A review of methodological evolution, theoretical foundations, and interdisciplinary frontiers toward next-generation scientific computing,” *Applied Sciences*, vol. 15, no. 14, p. 8092, 2025.

- [22] P. Paris and F. Erdogan, “A critical analysis of crack propagation laws,” *Journal of Basic Engineering*, vol. 85, no. 4, pp. 528–533, 1963.
- [23] X. Chen, M. Ma, Z. Zhao, Z. Zhai, and Z. Mao, “Physics-informed deep neural network for bearing prognosis with multisensory signals,” *Journal of Dynamics, Monitoring and Diagnostics*, pp. 200–207, 2022.
- [24] M. Parziale, L. Lomazzi, M. Giglio, and F. Cadini, “Physics-informed neural networks for the condition monitoring of rotating shafts,” *Sensors*, vol. 24, no. 1, p. 207, 2023.
- [25] N. Herwig, P. Borghesani, W. Smith, and Z. Peng, “Signal processing- and physics-informed neural network for explainable bearing condition monitoring,” *Mechanical Systems and Signal Processing*, vol. 235, p. 112925, 2025.
- [26] J. Zhong, Y. Zheng, C. Ruan, L. Chen, X. Bao, and L. Lyu, “M-IPISincNet: An explainable multi-source physics-informed neural network based on improved SincNet for rolling bearings fault diagnosis,” *Information Fusion*, vol. 115, p. 102761, 2025.
- [27] T. von Hahn and C. K. Mechefske, “Knowledge informed machine learning using a Weibull-based loss function,” *arXiv preprint arXiv:2201.01769*, 2022.
- [28] L. Lu, X. Meng, Z. Mao, and G. E. Karniadakis, “DeepXDE: A deep learning library for solving differential equations,” *SIAM Review*, vol. 63, no. 1, pp. 208–228, 2021.
- [29] A. Saxena, J. Celaya, E. Balaban, K. Goebel, B. Saha, S. Saha, and M. Schwabacher, “Metrics for evaluating performance of prognostic techniques,” in *2008 International Conference on Prognostics and Health Management*, 2008, pp. 1–17.

Chapter A

Reproducibility Notes

A.1 FINAL RESULT TABLE

The main result table is repeated here for quick reference.

Table A.1: Final prediction metrics for all models and experiments.

Experiment	Model	MAE	RMSE	R2
Exp 1	Data-only FNN	0.083	0.128	0.801
Exp 1	Proposed PINN	0.147	0.173	0.638
Exp 1	LSTM	0.101	0.112	0.848
Exp 1	CNN	0.244	0.266	0.146
Exp 2	Data-only FNN	0.104	0.148	0.738
Exp 2	Proposed PINN	0.114	0.144	0.751
Exp 2	LSTM	0.141	0.155	0.708
Exp 2	CNN	0.200	0.235	0.334
Exp 3	Data-only FNN	0.156	0.280	0.032
Exp 3	Proposed PINN	0.262	0.351	-0.530
Exp 3	LSTM	0.176	0.219	0.382
Exp 3	CNN	0.183	0.204	0.463

A.2 COMMANDS USED FOR THE LOCAL PIPELINE

```
uv sync --extra dev
uv run thesis-work validate-data
uv run thesis-work extract-features
uv run thesis-work run
uv run thesis-work regenerate-figures
uv run pytest -q
```

A.3 RAW DATA LOCATION

The raw IMS folders are expected in:

```
data/raw/  
  1st_test/  
  2nd_test/  
  3rd_test/
```

These folders are intentionally not part of the LaTeX source because they contain the local raw dataset.

A.4 LATEX BUILD

The LaTeX source is stored under `thesis/latex`. The file `build_pdf.bat` compiles the document into `physics_informed_bearing_rul_thesis.pdf` using MiKTeX.